

Chapter 5

Introduction to Trees

The Main Topics


- Learn about binary trees and the basic terminologies used in binary trees
- Explore various binary tree traversal algorithms
- Learn about
 - Binary search trees
 - Binary search balanced trees
 - Red-Black trees
 - B-Treesand explore how to implement the basic operations on them

Definition of a Tree

A *tree structure* (or *tree* for short) with base type T is either:

- The empty structure, or
- A node of type T , called *root*, with a finite number of associated disjoint tree structures with base type T , called *subtrees*

Representation of a Tree



Basic Concepts

- The top node is commonly called the *root*
- If an edge is between node n and node m , and node n is above node m in the tree:
 - n is the *parent* of m
 - m is a *child* of n
- ➡ The *parent-child* relationship between the nodes is generalized to the *ancestor-descendant* relationship
- A node that has no children is called a *leaf* of the tree
- A node which is not leaf is an *interior node*

Basic Concepts


- Children of the same parent are called *siblings*
- A *subtree* in a tree is any node in the tree together with all of its descendants
 - A subtree of a node n is a subtree rooted at a child of n
- The number of children of an interior node is called its *degree*
- The maximum degree over all nodes is the *degree of the tree*
 - ➡ The tree name is named after the degree of the tree

Basic Concepts

- The sequence of edges that connects an ancestor and a descendant is called a *path*
- The *height of a tree* is the number of nodes on the longest path from the root to a leaf
 - Alternatively, the number of edges of the longest path denotes the height
- If node n is at *level i* , then its child is said to be at *level $i + 1$* . The root of tree is at *level 1*
 - Some authors define the root is at *level 0*

Binary Trees

A *binary tree* is a finite set of nodes which either is empty or consists of a root with two disjoint binary trees called the *left* and the *right subtree* of the root



Representation of Binary Trees

- A binary tree consists of nodes of a type defined as follows:

```
typedef struct Node * Ref;  
struct Node {  
    int key;  
    Ref left, right;  
};
```

- An *empty* binary tree is represented as follows:

```
Ref root = NULL;
```

Constructing a Perfectly Balanced Binary Tree


- A binary tree is *perfectly balanced* if for each node the numbers of nodes in its left and right subtrees differ by at most 1
- An algorithm for constructing a PBBT with n nodes is best formulated in recursive terms:
 1. Use one node for the root
 2. Construct the left subtree with $n_l = \left\lfloor \frac{n}{2} \right\rfloor$ nodes *in this way*
 3. Construct the right subtree with $n_r = n - n_l - 1$ nodes *in this way*

Pseudo-code

```
Ref tree(n) {  
    if (n == 0)    return NULL;  
    n1 = n / 2;  
    nr = n - n1 - 1;  
    cin >> k;  
    Ref r = new Node;  
    r->key = k;  
    r->left = tree(n1);  
    r->right = tree(nr);  
    return r;  
}  
  
Ref root = tree(n);
```

Binary Tree Traversal

- A traversal algorithm for a binary tree visits each and every node in the tree
 - Assume that visiting a node simply means displaying the key portion of the node
- According to the recursive definition of a binary tree, the binary tree T is either empty or is of the form:



Binary Tree Traversal

- If T is empty, the traversal algorithm takes no action
 - An empty tree is the base case
- Otherwise, the traversal algorithm must perform three tasks: display the key in the root r , and traverse two subtrees T_L and T_R
- The algorithm has three choices when to visit r :
 - $r \rightarrow T_L \rightarrow T_R$ (*preorder traversal*)
 - $T_L \rightarrow r \rightarrow T_R$ (*inorder traversal*)
 - $T_L \rightarrow T_R \rightarrow r$ (*postorder traversal*)

Binary Tree Traversal: Pseudo-code

```
void preOrder(Ref r) {  
    if (r) {  
        cout << r->key;  
        preOrder(r->left);  
        preOrder(r->right);  
    }  
}  
  
void inOrder(Ref r) {  
    ...  
}  
  
void postOrder(Ref r) {  
    ...  
}
```

Binary Search Trees

- A *binary search tree* (or BST) is one that satisfies the following property:

“For each node its value is greater than all values in its left subtree and less than all values in its right subtree”
- BSTs are also called *ordered binary trees*

Binary Search Tree Traversal

- The traversal of a BST is the same as that of a binary tree
- An *inorder traversal* of a BST will visit the tree's nodes in sorted order according to their keys


Search Operation

- Firstly, check whether the root (of the whole tree or the current subtree) contains the search key k
 - If it does then the search process finishes successfully
 - If k is less than the value stored in the root, the search process continues with the left subtree
 - If k is greater than the value stored in the root, the search process continues with the right subtree
- If k is not in the tree, the search process always ends at an *empty* subtree

Search Operation: Pseudo-code

```
Ref  search(Ref r, int k) {  
    while (r)  
        if (r->key == k)  
            return r;  
        else  
            if (r->key > k)  
                r = r->left;  
            else  
                r = r->right;  
    return  NULL;  
}
```

Insertion Operation

- A new node is always inserted to a tree as a new leaf
 - It cannot be inserted to a node whose both subtrees are nonempty
 - Since BSTs are ordered ones, so the new key has to be inserted to an appropriate place in the tree
 - The algorithm must search for the right place first
-  The operation should be called *tree search with insertion* task

Insertion Operation: Rough Algorithm

- Firstly, the key will be searched for in the tree
- If the key is found, the algorithm does nothing more than returns back
- Otherwise, the search process will lead us to an empty subtree
- The new node must be inserted at the place of the empty tree

Insertion Operation: Pseudo-code

```
void searchAdd(Ref & r, int k) {  
    if (r == NULL) {  
        r = new Node;  
        r->key = k;  
        r->left = r->right = NULL;  
    }  
    else  
        if (r->key > k) searchAdd(r->left, k);  
        else  
            if (r->key < k) searchAdd(r->right, k);  
            else  
                return;  
}
```

Deletion Operation

- Similar to the insertion operation, this operation consists of search task and deletion task
 - ➡ This operation should be called *tree search with deletion* task
- Firstly, the key will be searched for in the tree
 - If the key is not found: Stop!!!
 - Otherwise, the search process will lead us to the node that contains the key
- ➡ It's time to run deletion task

Deletion Operation

- It is simple if the node to be deleted is a leaf or one with a single child
- The difficulty lies in removing a node with two children
 - A single pointer cannot point in two directions

Solution: The deleted node is to be replaced by either its *predecessor* or *successor*

- One of them will be actually deleted

Deletion Operation: Pseudo-code

```
void searchDel(Ref & r, int k) {  
    if (r == NULL) return;  
    if (r->key > k) searchDel(r->left, k);  
    else  
        if (r->key < k) searchDel(r->right, k);  
    else {          // ==  
        q = r;  
        if (q->right == NULL) r = q->left;  
        else  
            if (q->left == NULL) r = q->right;  
            else  
                del(r->left, q);  
        delete q;  
    }  
}
```


Deletion Operation: Pseudo-code

```
void del(Ref & r, Ref & q) {  
    if (r->right)  
        del(r->right, q);  
    else {  
        q->key = r->key;  
        q = r;  
        r = r->left;  
    }  
}
```

```
void searchDel(Ref & r, int k)  
{  
    ...  
    else { // ==  
        q = r;  
        if (q->right == NULL)  
            r = q->left;  
        else  
            if (q->left == NULL)  
                r = q->right;  
            else  
                del(r->left, q);  
        delete q;  
    }  
}
```

Balanced Binary Search Trees

- In 1962, two Soviet scientists, Adelson-Velskii and Landis, proposed a new type of BSTs called *balanced binary search tree*
 - Nowadays, it's well known as AVL tree – named after its inventors
- The noteworthy point of AVL trees is its balance criterion

The Balance Criterion

- The criterion is stated as follows:



A binary tree is balanced if and only if for every node the heights of its two subtrees differ by at most 1

- The criterion makes AVL trees have 3 advantages over the others:
 - The height of AVL trees is at most $1.44 \log_2 n$
 - AVL trees could not be deformed as BSTs
 - The cost of rebalancing AVL trees is nearly constant time and in the worst case, it's $O(\log n)$

The Basic Operations

- The search algorithm for an AVL tree is the same as the search algorithm for a BST
- Insertion and deletion operations on AVL trees are somewhat different from the ones discussed for BSTs
 - After inserting a node to (or deleting a node from) an AVL tree, the resulting tree may still be an AVL tree or a *rebalancing* must be done to restore the balance criterion
- Three operations run in $O(\log n)$

When Does an Imbalance Arise?




- Given a root r with the left and right subtrees: T_L and T_R
 - In general, r can be any node in an AVL tree
- Let h_L and h_R be the heights of T_L and T_R , respectively
- Assume that the new node is inserted to T_L causing its height to increase by 1

When Does an Imbalance Arise?


Before	After
$h_L = h_R$	$h_L > h_R$: h_L and h_R become of unequal height, but the balance criterion is not violated
$h_L < h_R$	$h_L = h_R$: h_L and h_R obtain equal height, the balance has even been improved
$h_L > h_R$	The balance criterion is violated, and the tree must be restructured

The Rebalancing Procedure



- It's also called *rotating* the tree
- Suppose the rotation occurs at a node r :
 - *Left rotation*: Certain nodes from the right subtree of r move to its left subtree; the root of the right subtree of r becomes the new root of the reconstructed subtree

The Rebalancing Procedure



- It's also called *rotating* the tree
- Suppose the rotation occurs at a node r :
 - *Left rotation*: Certain nodes from the right subtree of r move to its left subtree; the root of the right subtree of r becomes the new root of the reconstructed subtree
 - *Right rotation*: Certain nodes from the left subtree of r move to its right subtree; the root of the left subtree of r becomes the new root of the reconstructed subtree

The Rebalancing Procedure: Case 1

The Rebalancing Procedure: Case 2

The Rebalancing Procedure: Case 3

The Rebalancing Procedure: Case 4

Insertion Operation

- Assume that the key to be added to the AVL tree is the new one
- Insertion operation includes 2 stages:
 1. Search the tree and add the new node to the appropriate place
 2. After inserting the new node in the tree, the resulting tree might not be an AVL tree
 - ➡ The rebalancing procedure must be activated

Insertion Operation

- How does the rebalancing procedure work?
 - It's accomplished by retreating along the search path and check if the balance criterion is violated at each node
 - If the balance criterion is violated at a node, the subtree rooted at that node will be rebalanced
- Once balance is established, the subtree no longer grew in height
 - ➡ The algorithm can ignore the remaining nodes on the path back to the root

Deletion Operation

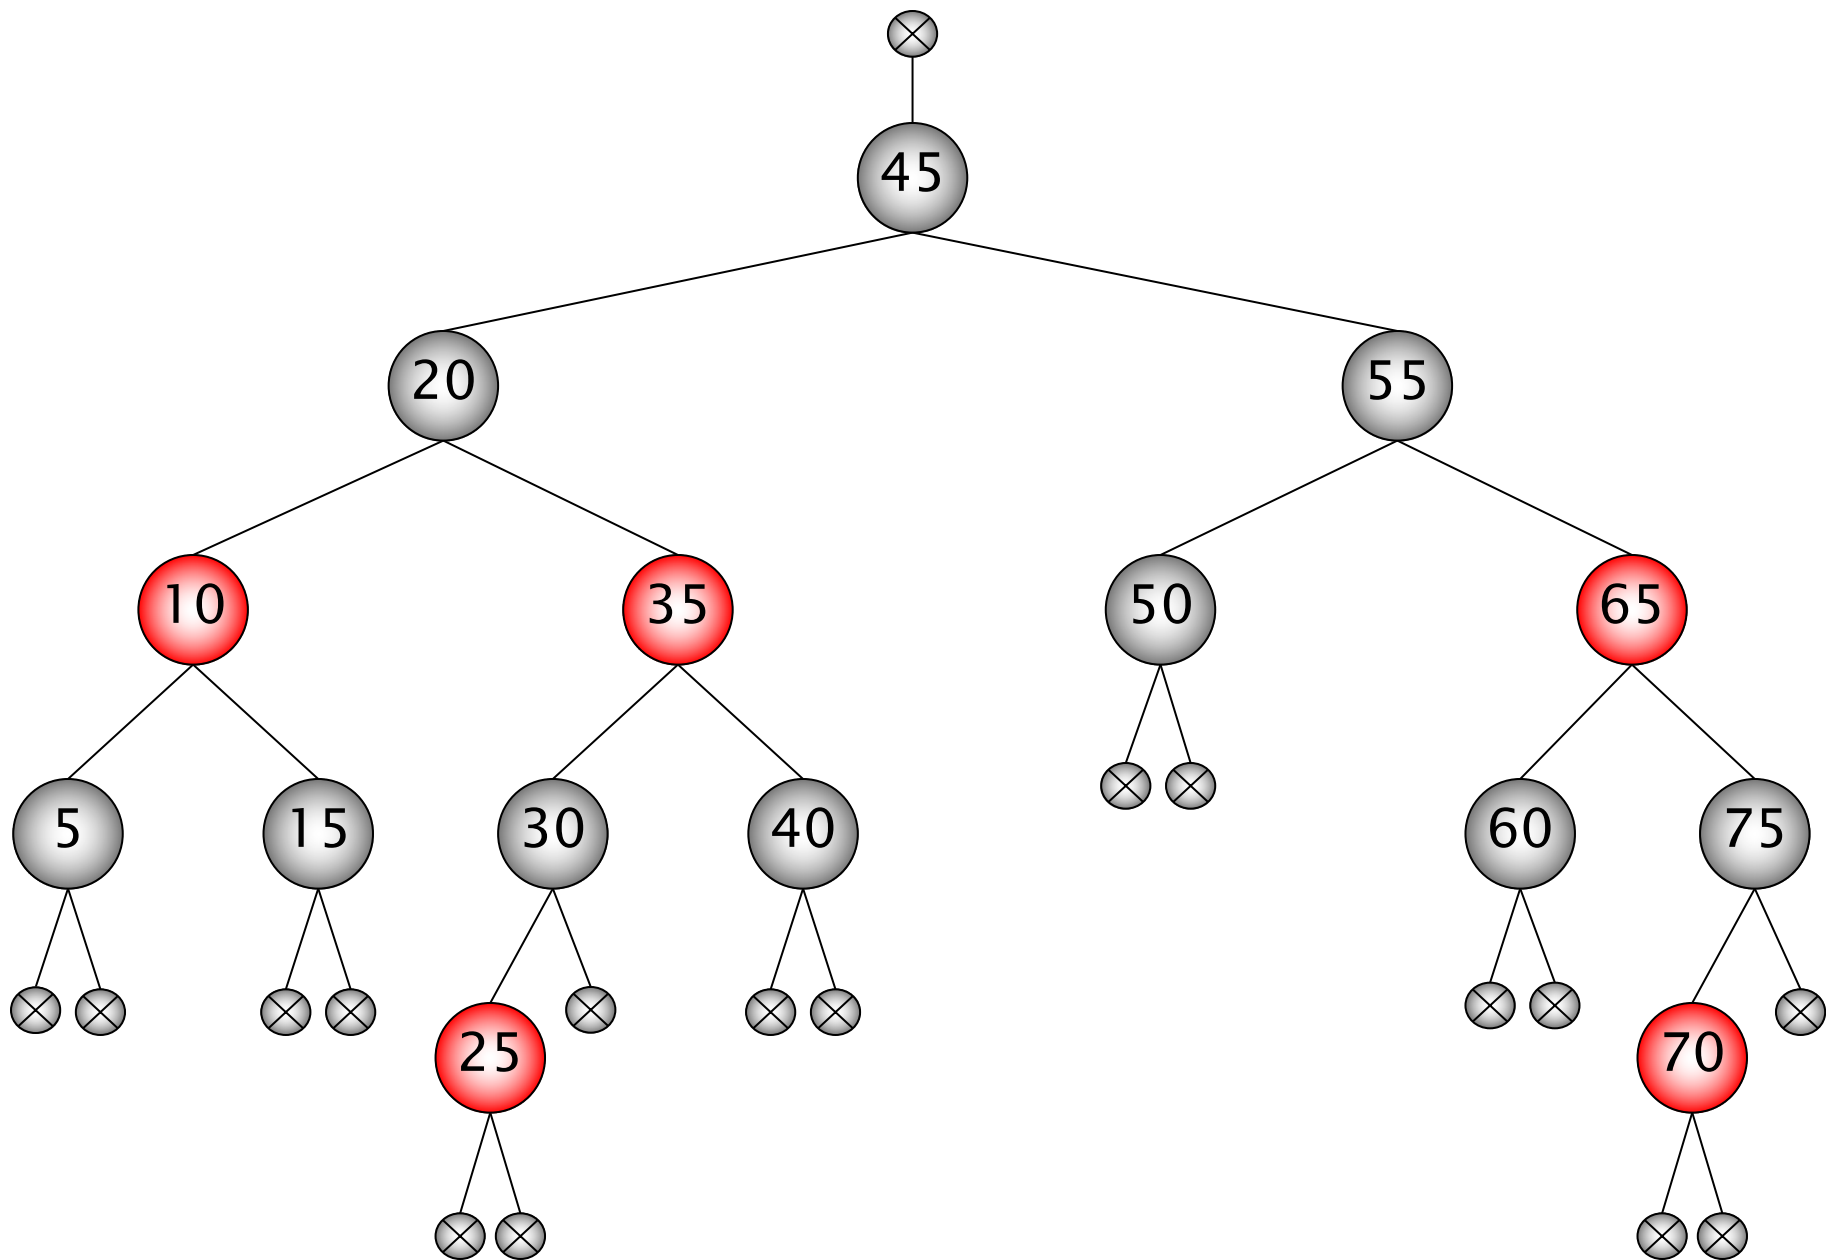
- Assume that the key to be deleted is in the AVL tree
- Deletion operation includes 2 stages:
 1. Find the node containing the key to be deleted and delete it
 2. After deleting the node in the tree, the resulting tree might not be an AVL tree
 - ➡ The rebalancing procedure must be activated

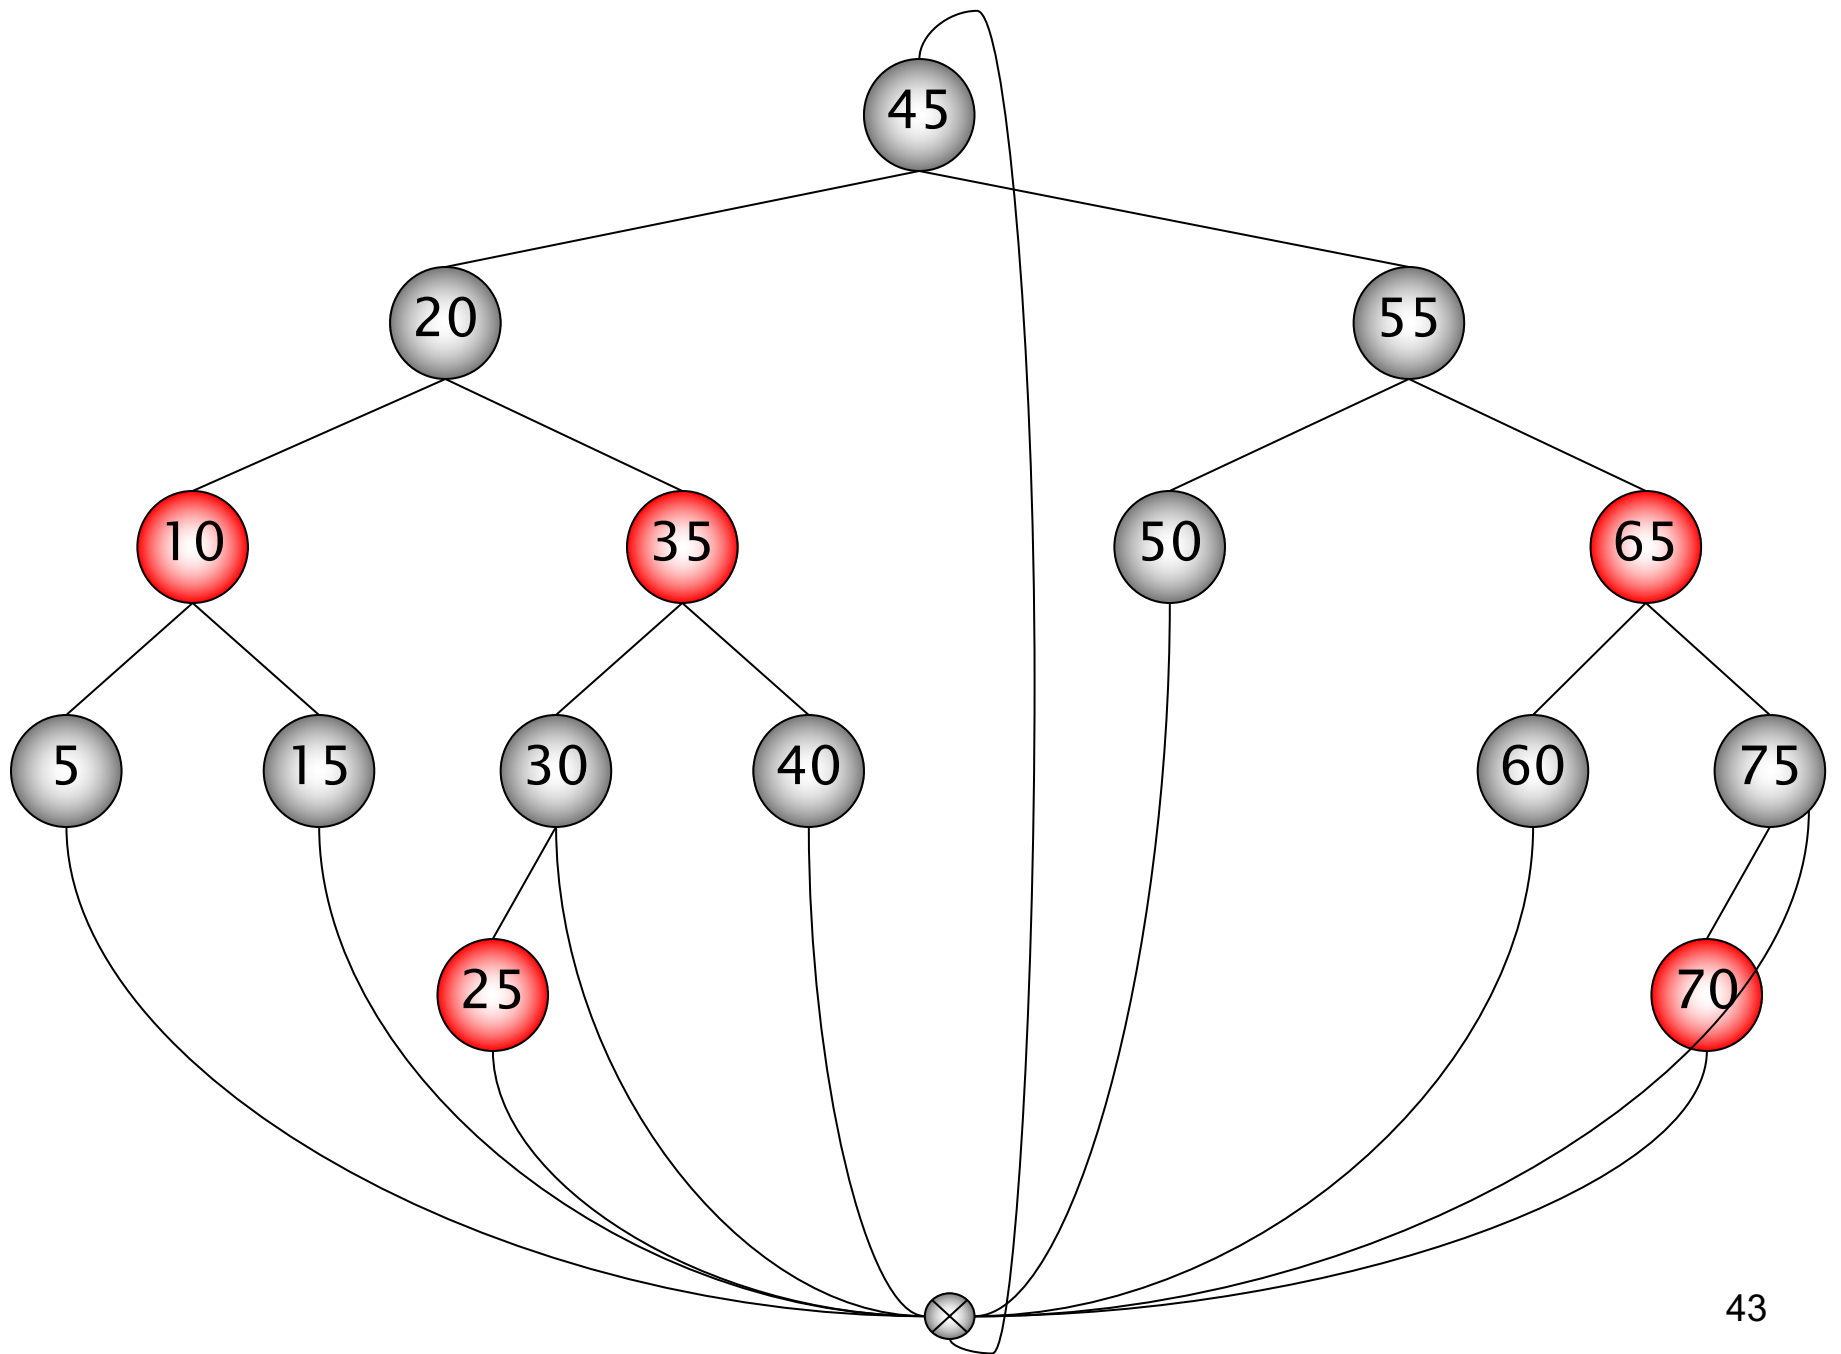
Deletion Operation

- How does the rebalancing procedure work?
- Its performance is similar to the case of insertion operation except two things:
 - After rebalancing a subtree, the overall tree might still not be an AVL tree
 - ➡ The algorithm has to continuously traverse back to the root node
 - The second thing?

Red-Black Trees

- A *red-black tree* (RBT) is a type of self-balancing BST
 - Each node in an RBT is labeled as *red* or *black*
- Operations on RBTs take $O(\log n)$ time in the worst case since the height of an RBT is at most $2 \log_2(n + 1)$
 - The height of an AVL tree is at most $1.44 \log_2 n$
- However, a careful nonrecursive implementation can be done relatively effortlessly compared with AVL trees





Definition of Red-Black Trees

An RBT is a BST that satisfies the following *red-black criteria*:

1. Every node is either *red* or *black*
2. The root of the tree is always *black*
3. All leaves are *black*
4. If a node is *red*, then its parent is *black*
5. Any path from a node to any of its leaves contains the same number of black nodes, called *black height*

Representation of Red-Black Trees

```
typedef struct Node * Ref;
struct Node {
    int key;
    int color;
    Ref parent, left, right;
};

ref getNode(int key, int color, Ref nil) {
    p = new Node;
    p->key    = key;
    p->color  = color;
    p->left   = p->right = p->parent = nil;
    return  p;
}
```

The Initial State of a Red-Black Tree

```
Ref nil, root;
```

```
...
```

```
nil = new Node;
```

```
nil->color = BLACK;
```

```
nil->key = -1;
```


```
nil->left =
```

```
nil->right =
```


```
nil->parent = nil;
```

```
...
```


```
root = nil;
```



Tree Rotation



Tree Rotation




```
leftRotate(Ref & root, Ref x) {
```

```
    y = x->right;
```

```
    x->right = y->left;
```

```
    if (y->left != nil)
```

```
        y->left->parent = x;
```

```
    y->parent = x->parent;
```

```
    if (x->parent == nil) root = y;
```

```
    else
```

```
        if (x == x->parent->left)
```

```
            x->parent->left = y;
```



```
        else
```

```
            x->parent->right = y;
```

```
    y->left = x;
```

```
    x->parent = y;
```

```
}
```



Search Operation

- An RBT is a BST so the search algorithm for an RBT is the same as the search algorithm for a BST

Insertion Operation

- The new node, as usual, is placed as a leaf in the tree
- The node must be colored *red*
 - If the parent is *black*: The red-black criteria are maintained
 - If the parent is *red*: The criterion “no two consecutive red nodes” is violated
 - ➡ Rebalancing the tree is needed

Insertion Operation: Pseudo-code


```
void    RBT_Insertion(Ref & root, int key) {  
    x = getNode(key, RED, nil);  
    BST_Insert(root, x);  
    Insertion_FixUp(root, x);  
}
```

```

BST_Insert(Ref & root, ref x) {
    y = nil;
    z = root;
    while (z != nil) {
        y = z;
        if (x->key < z->key)          z = z->left;
        else if (x->key > z->key)      z = z->right;
        else                          return;
    }
    x->parent = y;
    if (y == nil)    root = x;
    else
        if (x->key < y->key)  y->left = x;
        else                y->right = x;
}


```

Some Conventions




- x : The pointer designated to point to the newly added leaf
- p : The pointer designated to point to the *parent* of the node pointed to by x
- u : The pointer designated to point to the *uncle* of the node pointed to by x
- g : The pointer designated to point to the *grandparent* of the node pointed to by x


Imbalance Cases



Case 1



Case 2




Case 3

The Strategies for Rebalancing an RBT

- Case 1


→ Reverse the color of three nodes: u , p , and g



- Case 2

→ Reverse the color of two nodes: p and g


→ Run a rotation at g




- Case 3

→ Run a rotation at parent p


→ Go to Case 2




Case 1




Reverse color




Case 2




Case 2




Case 3





Case 3



Right rotation



Reverse color



Left rotation

Insertion Operation: Pseudo-code

```
Insertion_FixUp(Ref & root, Ref x) {  
    while (x->parent->color == RED)  
        if (x->parent == x->parent->parent->left)  
            ins_leftAdjust(root, x);  
        else  
            ins_rightAdjust(root, x);  
    root->color = BLACK;  
}
```


```

ins_leftAdjust(Ref & root, Ref & x) {
    u = x->parent->parent->right;
    if (u->color == RED) {
        x->parent->color = u->color = BLACK;
        x->parent->parent->color = RED;
        x = x->parent->parent;
    }
    else {
        if (x == x->parent->right) {
            x = x->parent; leftRotate(root, x);
        }
        x->parent->color = BLACK;
        x->parent->parent->color = RED;
        rightRotate(root, x->parent->parent);
    }
}

```

B-Trees

Example: A B-tree of Order 2




Definition of B-Trees

A B-tree of order t is either empty, or has the following properties:

- Every node contains at most $2t$ keys
- All nodes, except the root, contain at least t keys
- Every node is either a leaf or it has $m + 1$ children, where m is its number of keys
- All leaves are on the same level

Example: A B-tree of Order 2



The Structure of a Node (or Page)

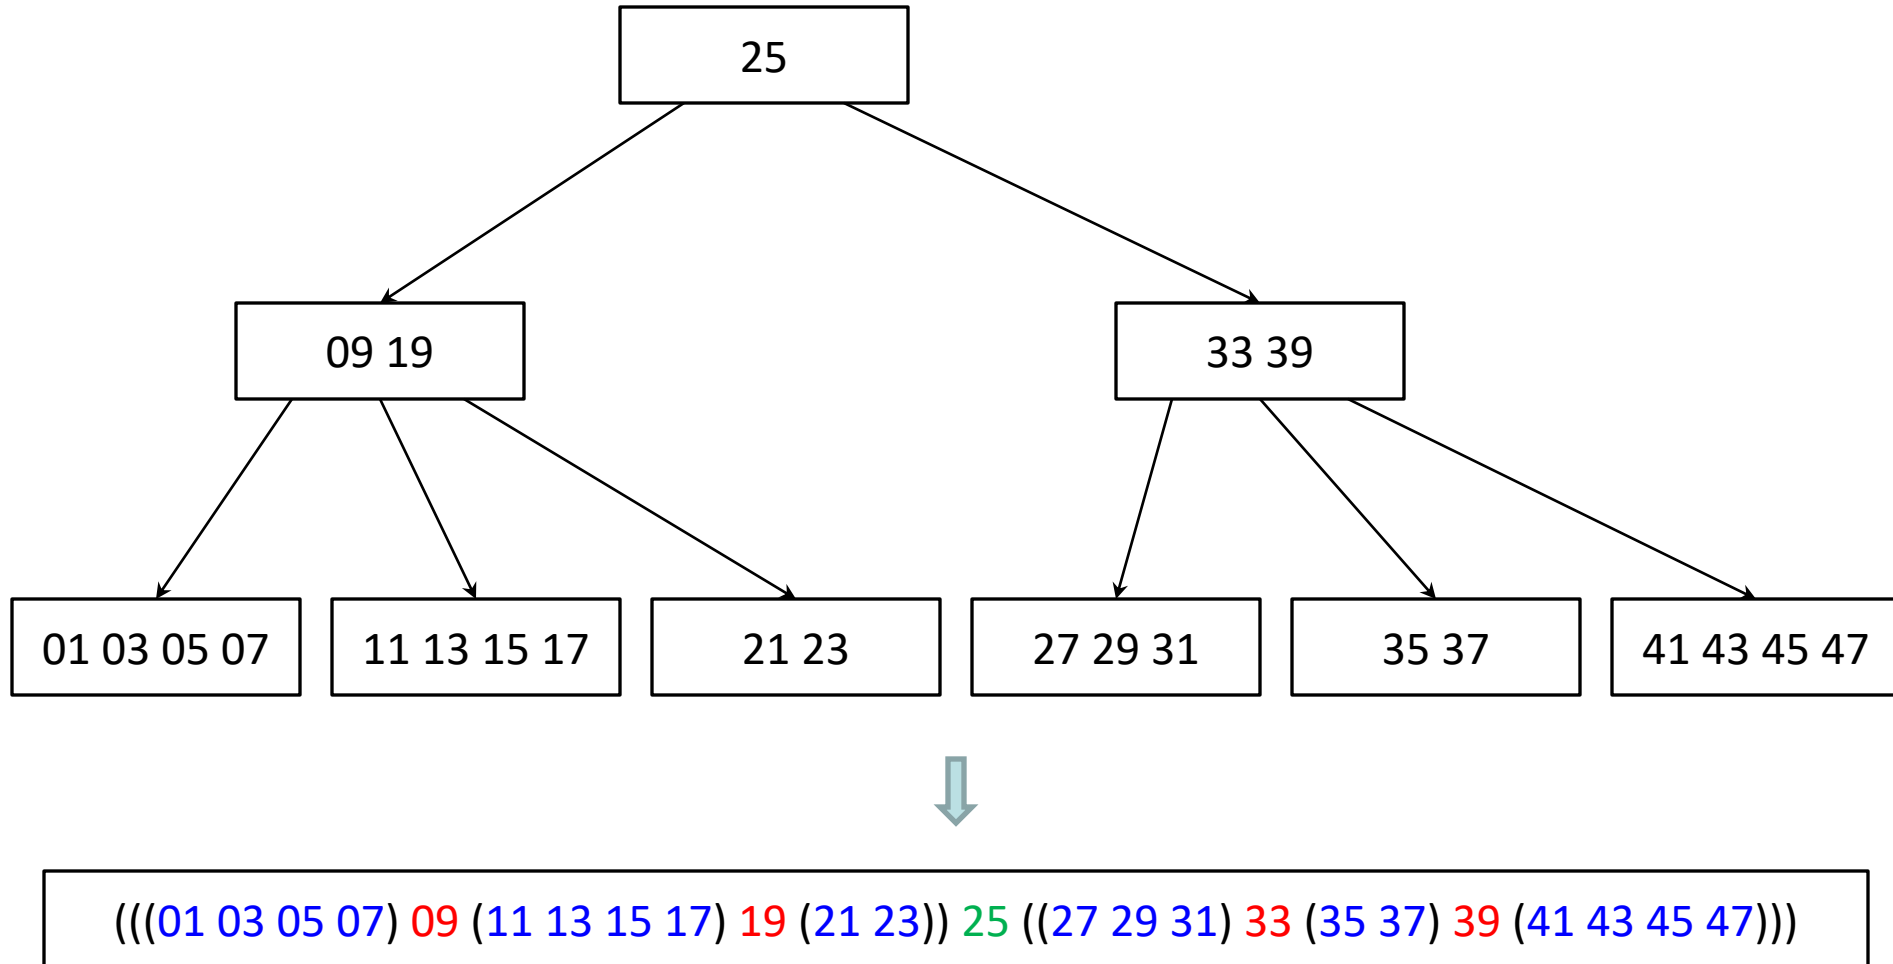
The structure of a node is as follows:

p_0	k_1	p_1	k_2	p_2	\dots	k_m	p_m
-------	-------	-------	-------	-------	---------	-------	-------

where

- $k_1 < k_2 < \dots < k_m$
- p_i is a pointer to a child
 - If it's a leaf: $p_i = \text{NULL}, \forall i \in [0, m]$
- All keys in the node to which p_i points are greater than k_i and less than k_{i+1}


Example



Search Operation

- The search must start at the root of tree
 - Assume that the node being considered contains m keys:
 k_1, k_2, \dots, k_m
 - If m is sufficiently large, one may use binary search; otherwise, a sequential search will do
- Let k be the search key. If the search is unsuccessful:
 - $k < k_1$: Search the node pointed to by p_0
 - $k > k_m$: Search the node pointed to by p_m
 - $k_i < k < k_{i+1}$: Search the node pointed to by p_i
- If the designated pointer is a null pointer: Stop!!!

Example



Insertion Operation

- Assume that the key to be inserted is new
 - ➡ The search process terminates at a leaf
- The new key is inserted into the leaf if there is room
- If the leaf is full
 - (for pedagogical reasons) *Insert the new key into the leaf*
 - Split the leaf into two nodes
 - Move the median key to the parent node
- The splitting can propagate upward up to the root, causing the tree to increase in height

Example

Deletion Operation

- Assume that the key to be deleted, say k , is in the tree
 - ➡ The search process terminates at the node containing k
- There are two different circumstances:
 - It's a leaf: The removal algorithm is plain and simple
 - Otherwise: The key must be replaced by its predecessor or successor, which happen to be on leaves
- ➡ In either cases, the key that actually to be deleted is always on a leaf

Deletion Operation

- If the leaf contains more than t keys
 - Delete k and no further action is required
 - If the leaf contains only t keys
 - If one of the adjacent siblings has more than t keys: Move one key from that sibling to the parent and one key from the parent to the leaf, and then delete k
 - Otherwise: Combine one of the adjacent siblings with the leaf and the median key from the parent, and then delete k
- ➡ This process may propagate all the way up to the root which could result in reducing the height of the B-tree

Example

B-Trees: Summary

- In practice, B-trees are designed to store and manage a large data on secondary storage devices
- A node of a B-tree normally corresponds to a *disk page*
 - For a typical disk, a page might be 2^{11} to 2^{14} bytes in length
- The time needed to access a disk page is typically $\sim 10^5$ larger than the time needed to compare keys in RAM
 - ➡ The height of B-trees is the principal indicator of the efficiency of this data structure